



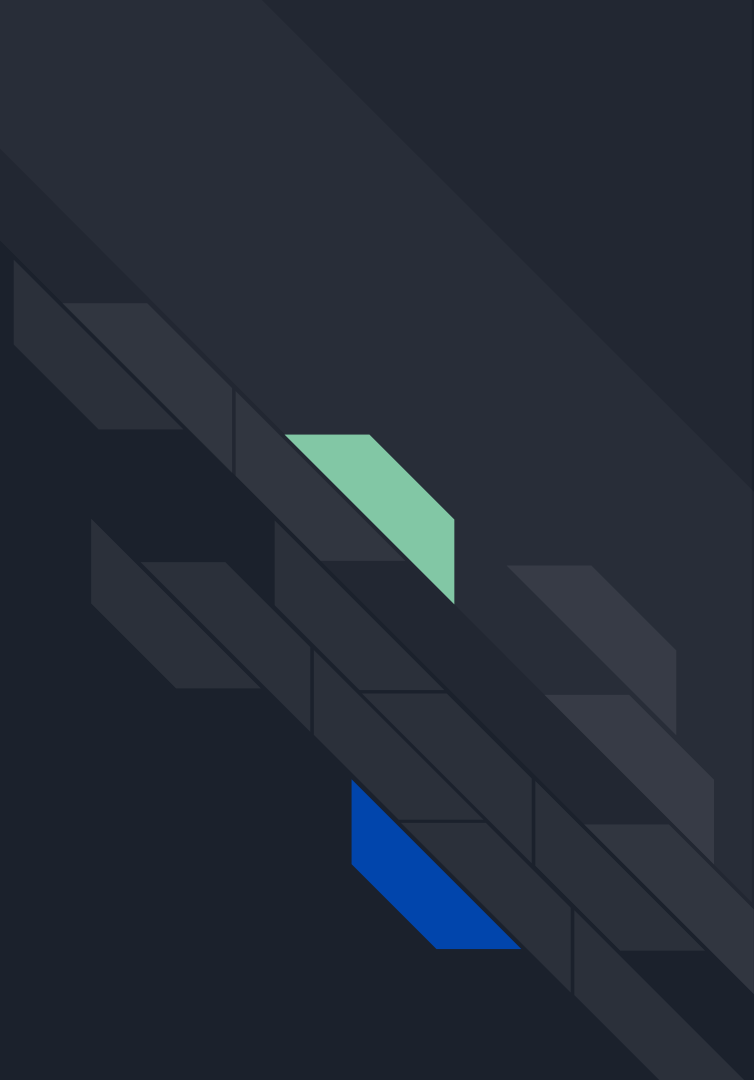
Teraslice

Stream Processing with NodeJS

<https://github.com/terascope/teraslice>

Austin Godber
@godber@az.social
May 17th, 2023

Should I use Teraslice?













Outline

- ETL and Stream Processing
- Teraslice Description and Features
- Teraslice Job
- Example Problem
- Teraslice Processor



Outline

- ETL and Stream Processing
- Teraslice Description and Features
- Teraslice Job
- Example Problem
- Teraslice Processor

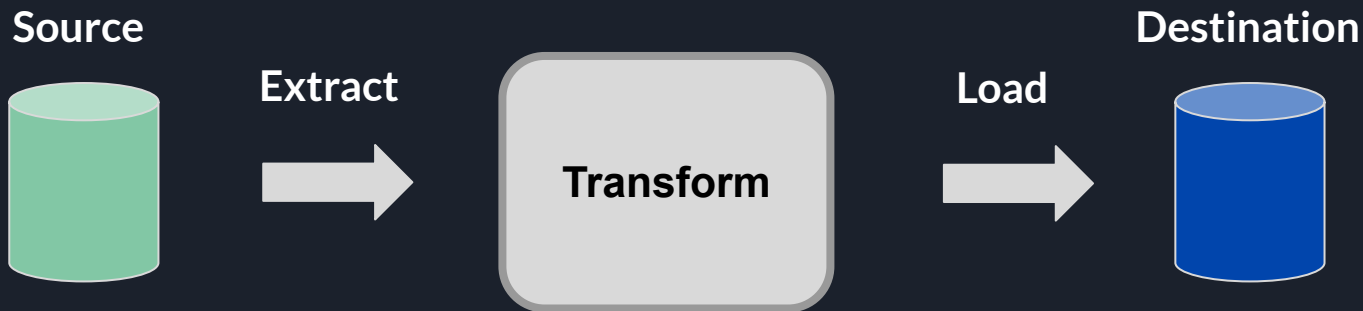


Extract Transform and Load

- Extract
 - Reading data from a data source like a database
 - Validation of the read data
- Transform
 - Modify the data according to need
 - Filtering
 - Mapping
 - Creating derived values
 - Deduplication
 - Enrichment from other sources
 - etc
- Load
 - Write modified data to destination system
 - Write audit log



Extract Transform and Load

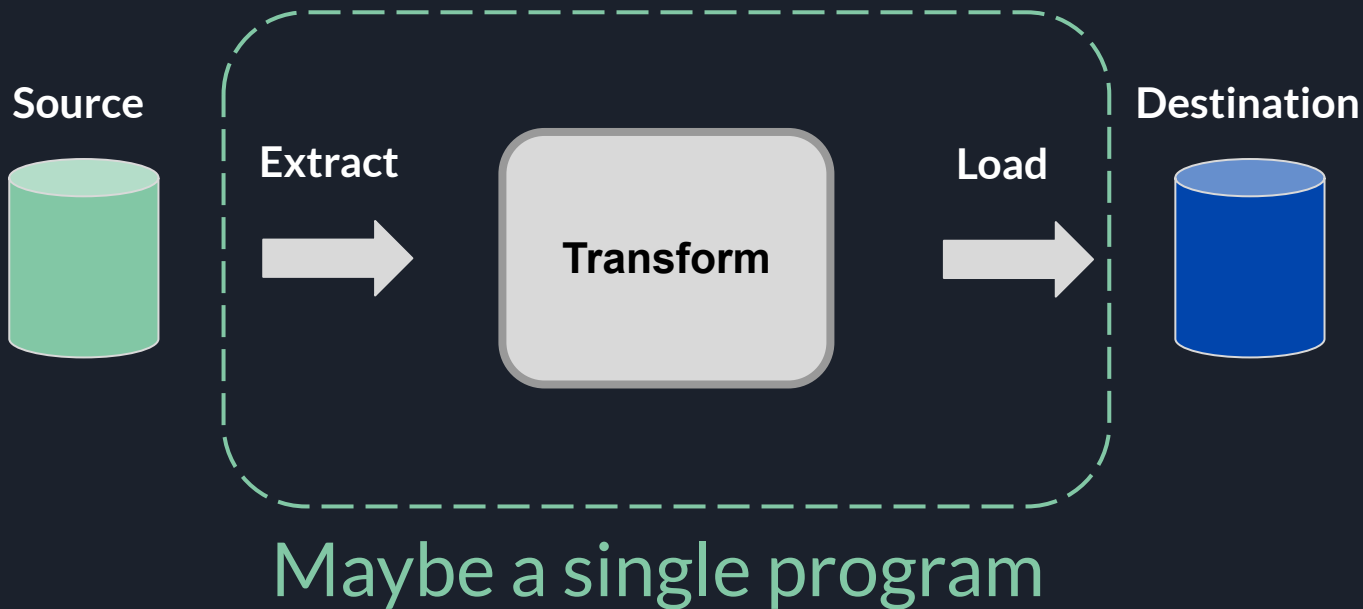


Extract Transform and Load

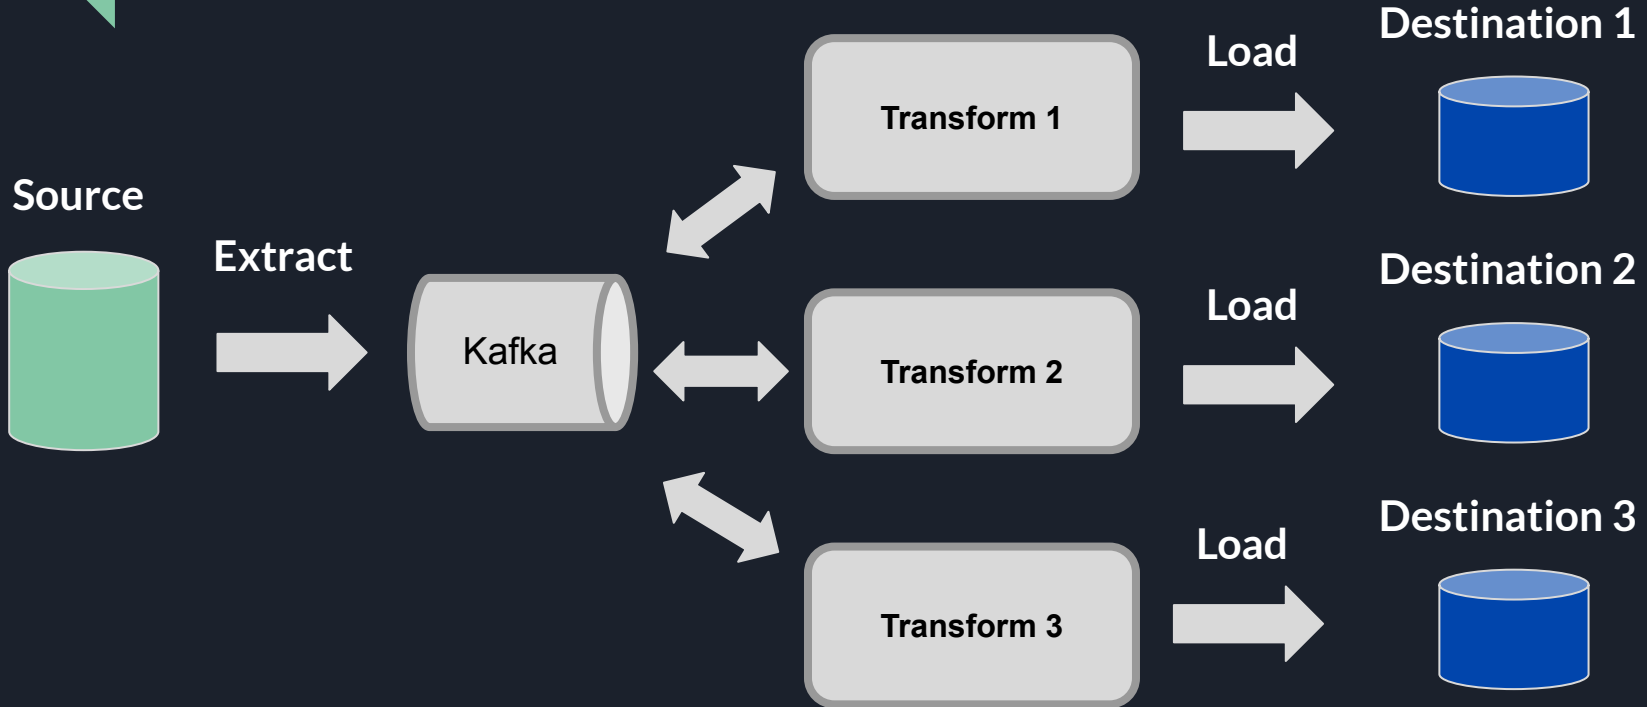


Usually batchy!

Extract Transform and Load



Streaming ETL





Outline

- ETL and Stream Processing
- Teraslice Description and Features
- Teraslice Job
- Example Problem
- Teraslice Processor



Teraslice Background

“Teraslice provides scalable data processing pipelines implemented using components written in JavaScript. It uses a distributed model to spread work across a cluster of computers and can easily process millions of records per second.”

<https://terascope.github.io/teraslice/>



Teraslice Background

Features

- Originally designed to “re-index” large (100B+) Elasticsearch Indices
- Enhanced to become a generic stream processing tool
- Written in Javascript/Typescript
- Custom “processors” can be implemented in Javascript and uploaded separately
- Management REST API
 - <https://terascope.github.io/teraslice/docs/management-apis/overview>
- Jobs are defined in JSON
 - <https://terascope.github.io/teraslice/docs/jobs/configuration#examples>
- Runs locally with native clustering mechanism or in Kubernetes



Teraslice Background

Components

- <https://github.com/terascopes/teraslice>
- <https://github.com/terascopes/elasticsearch-assets>
- <https://github.com/terascopes/kafka-assets>
- <https://github.com/terascopes/file-assets>
- <https://github.com/terascopes/standard-assets>



Teraslice Alternatives

- Vector (vector.dev) (Rust)
- Spark Streaming (Java)
- Flink (Java)
- Kafka Streams (Java)
- Opensearch Data Prepper Pipelines (Config)
- Elasticsearch Ingest Pipelines (Config)



Outline

- ETL and Stream Processing
- Teraslice Description and Features
- Teraslice Job
- Example Problem
- Teraslice Processor

```
{
  "name": "Data Generator To Kafka",
  "lifecycle": "persistent",
  "workers": 1,
  "assets": [
    "standard",
    "kafka"
  ],
  "operations": [
    {
      "_op": "data_generator",
      "size": 500000
    },
    {
      "_op": "kafka_sender",
      "connection": "default",
      "topic": "test1",
      "size": 25000
    }
  ]
}
```



```
{
  "name": "Data Generator To Kafka",
  "lifecycle": "persistent",
  "workers": 1,
  "assets": [
    "standard:v1.1.15",
    "kafka"
  ],
  "operations": [
    {
      "_op": "data_generator",
      "size": 500000
    },
    {
      "_op": "kafka_sender",
      "connection": "default",
      "topic": "test1",
      "size": 25000
    }
  ]
}
```

**Pin
Version**



```
{  
  "name": "Data Generator To Kafka",  
  "lifecycle": "persistent",  
  "workers": 1,  
  "assets": [  
    "standard",  
    "kafka"  
  ],  
  "operations": [  
    {  
      "_op": "data_generator",  
      "size": 500000  
    },  
    {  
      "_op": "kafka_sender",  
      "connection": "default",  
      "topic": "test1",  
      "size": 25000  
    }  
  ]  
}
```



Processors



```
{
  "name": "Data Generator To Kafka",
  "lifecycle": "persistent",
  "workers": 1,
  "assets": [
    "standard",
    "kafka"
  ],
  "operations": [
    {
      "_op": "data_generator",
      "size": 500000
    },
    {
      "_op": "kafka_sender",
      "connection": "default",
      "topic": "test1",
      "size": 25000
    }
  ]
}
```

Source

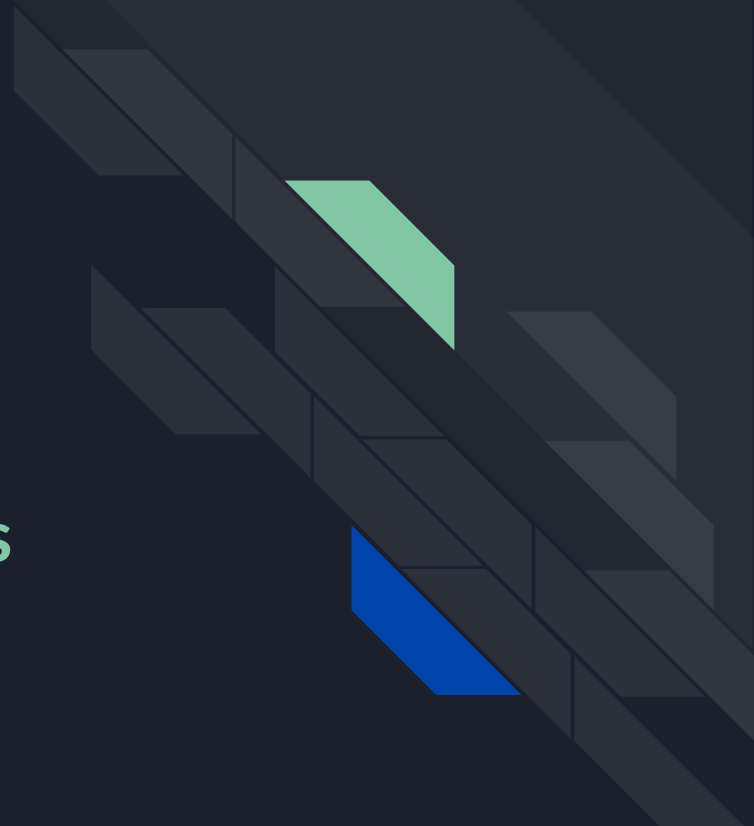
Destination



```
{  
  "name": "Data Generator To Kafka",  
  "lifecycle": "persistent",  
  "workers": 1,  
  "assets": [  
    "standard",  
    "kafka"  
  ],  
  "operations": [  
    {  
      "_op": "data_generator",  
      "size": 500000  
    },  
    {  
      "_op": "kafka_sender",  
      "connection": "default",  
      "topic": "test1",  
      "size": 25000  
    }  
  ]  
}
```



Insert Mutations & Transforms





Outline

- ETL and Stream Processing
- Teraslice Description and Features
- Teraslice Job
- Example Problem
- Teraslice Processor



Weather Station Streaming Problem

Imagine you operate thousands of weather stations all around the world and you had customers who want real time alerts from those weather stations under specified conditions.

Each measurement from one of these weather stations looks like this:

```
{  
  "station_id": "USW00003192",  
  "date": "2016-04-23T00:00:00",  
  "AWND": 2.4,  
  "PRCP": 0,  
  "TMAX": 29.4,  
  "TMIN": 18.9,  
}
```




Weather Station Streaming Problem

Imagine you operate thousands of weather stations all around the world and you had customers who want real time alerts from those weather stations under specified conditions.

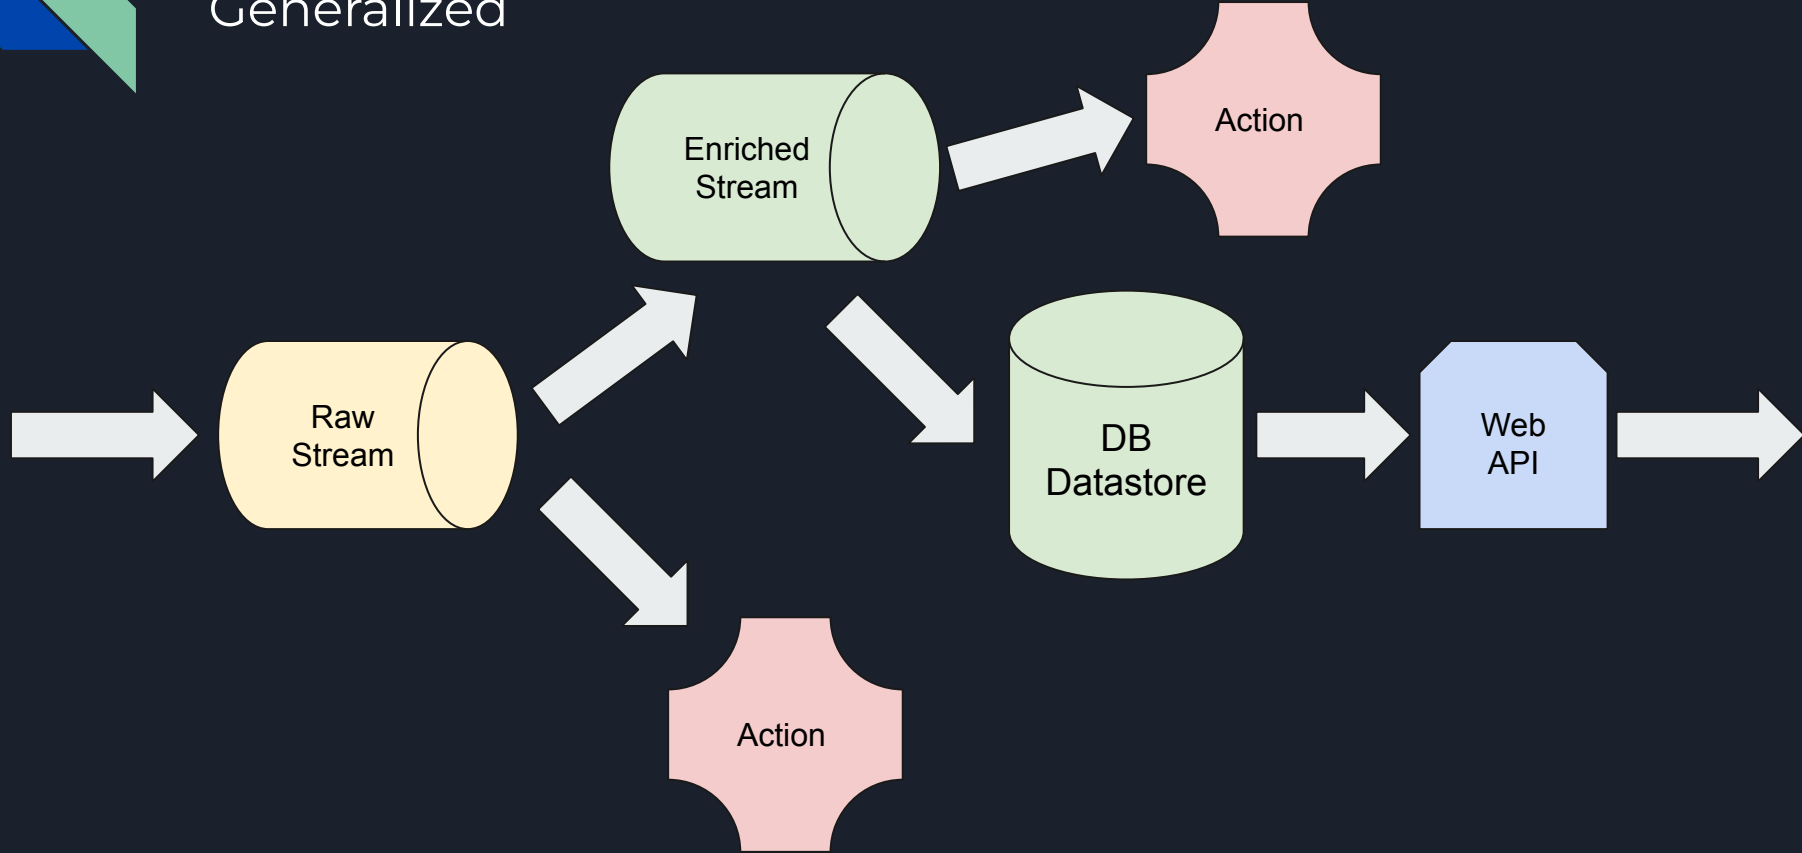
Each measurement from one of these weather stations looks like this:

Or more likely ...

```
"USW00003192", "2016-04-23T00:00:00", 2.4, 0, 29.4, 18.9
```

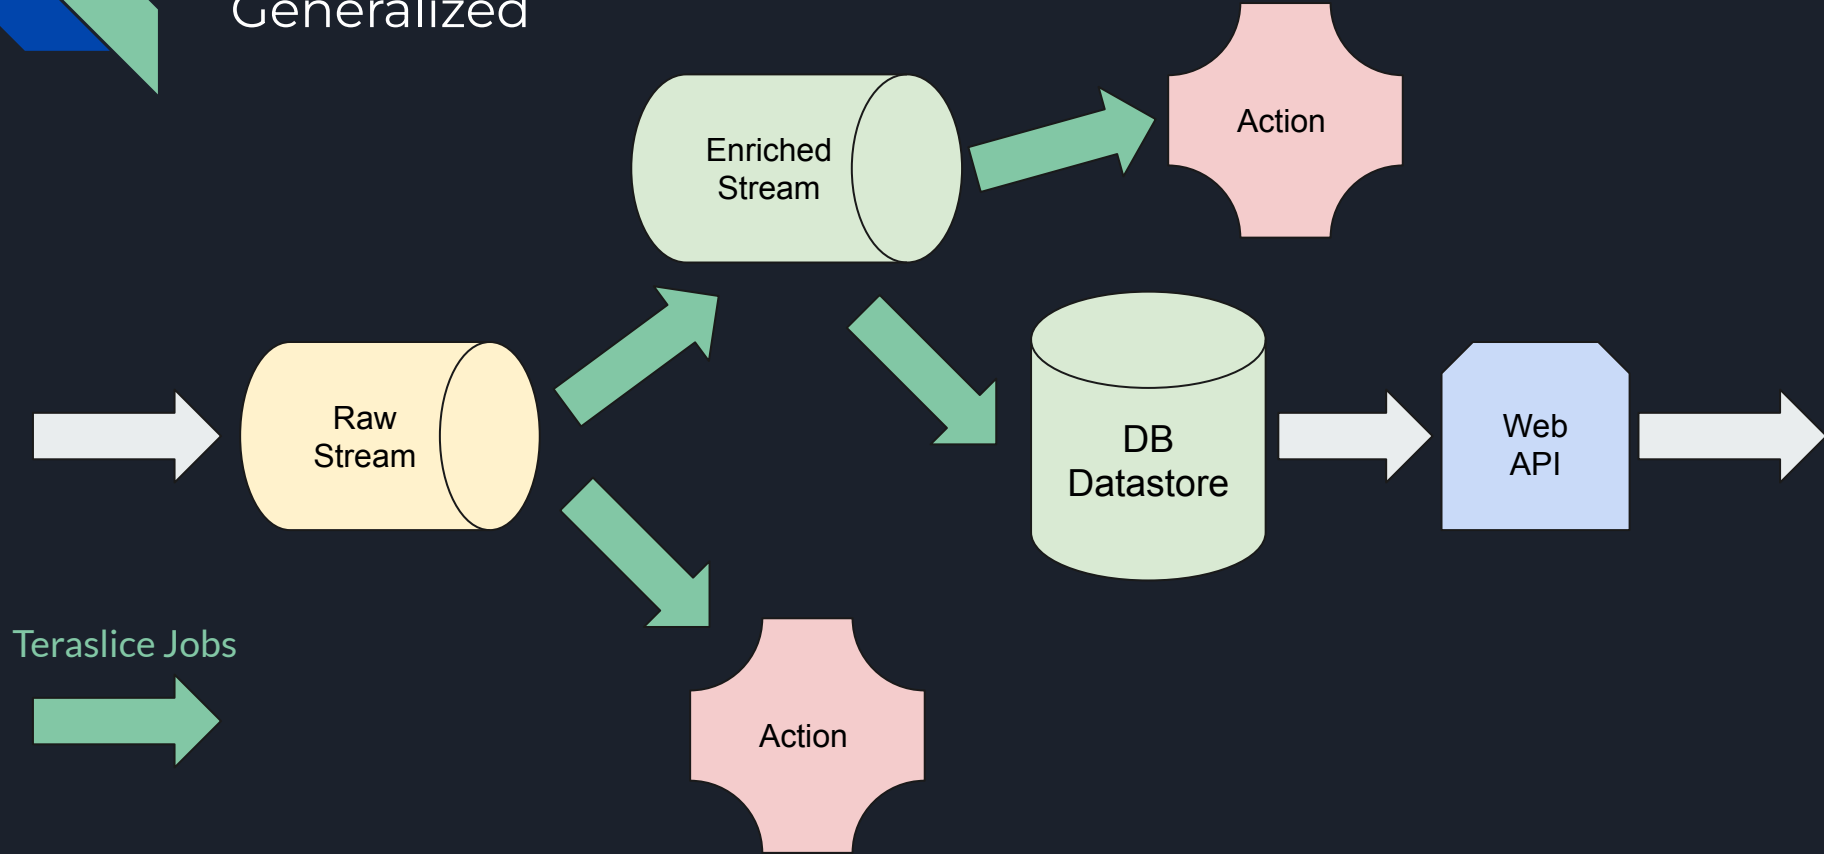
Weather Station Streaming Problem

Generalized



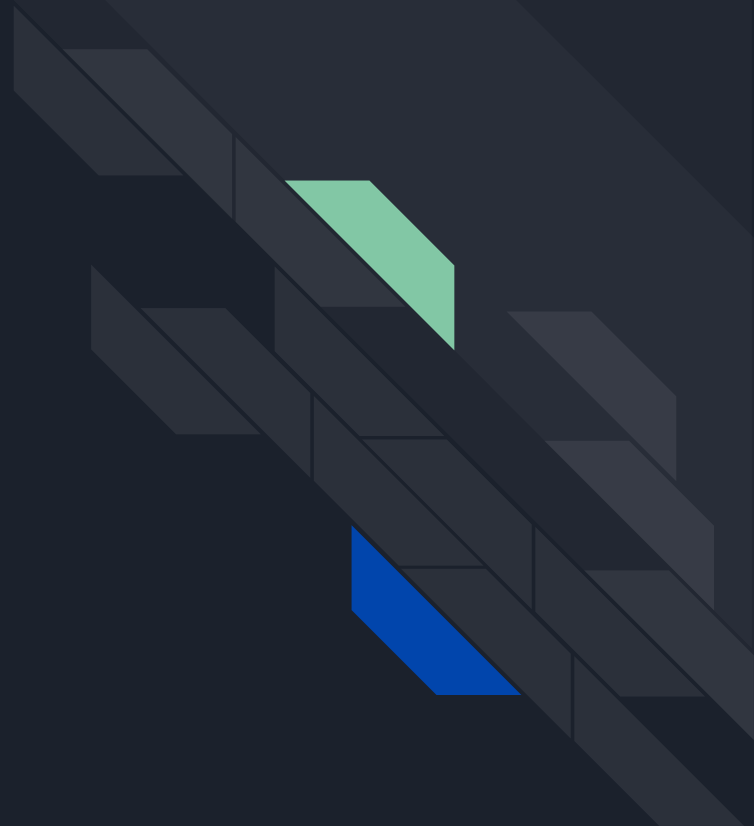
Weather Station Streaming Problem

Generalized



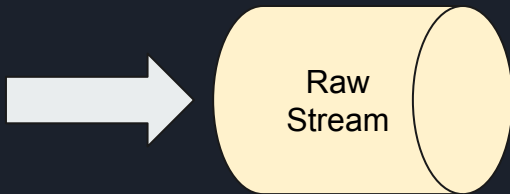
PAUSE

Batch based ETL workflows are **great** for many use cases, but as the incoming rate increases the batch processing time gets larger. When that processing time is **TOO** long for your use case, that's when you consider **switching** to streaming ... if you can.



Weather Station Streaming Problem

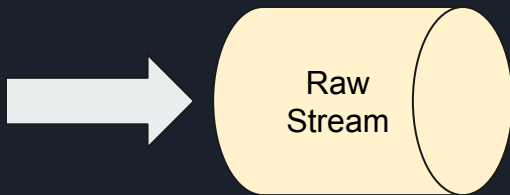
Incoming data is a
big array of JSON
records



```
{  
  "station_id": "USW00003192",  
  "date": "2016-04-23T00:00:00",  
  "AWND": 2.4,  
  "PRCP": 0,  
  "TMAX": 29.4,  
  "TMIN": 18.9,  
}
```

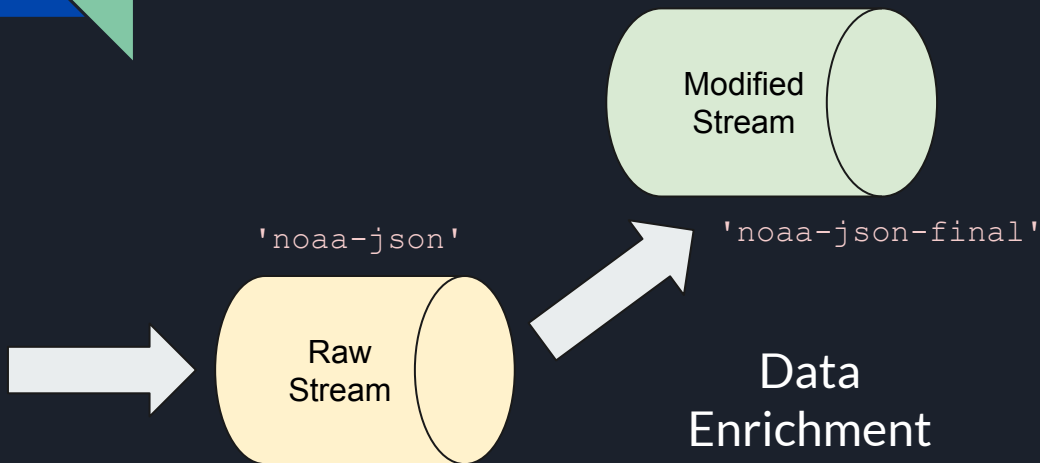
Weather Station Streaming Problem

Incoming data is
global and needs
interpretation



```
{  
  "station_id": "USW00003192",  
  "date": "2016-04-23T00:00:00",  
  "AWND": 2.4,  
  "PRCP": 0,  
  "TMAX": 29.4,  
  "TMIN": 18.9,  
}
```

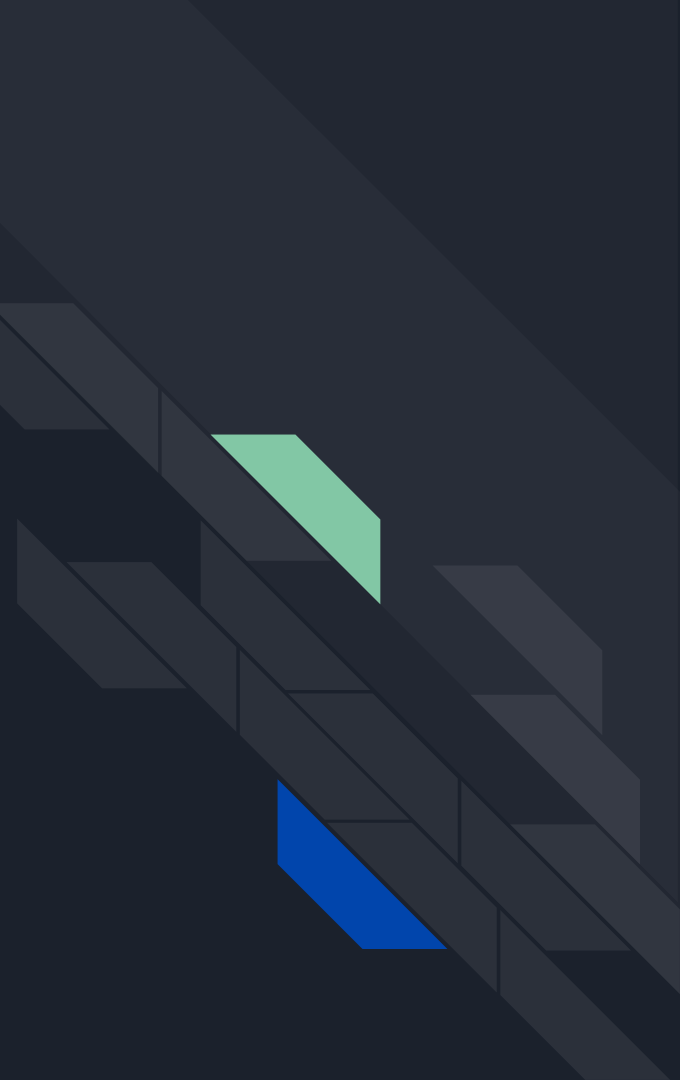

Weather Station Streaming Problem



```
{  
  "station_id": "USW00003192",  
  "date": "2016-04-23T00:00:00",  
  "AWND": 2.4,  
  "PRCP": 0,  
  "TMAX": 29.4,  
  "TMIN": 18.9,  
}
```

```
{  
  "station": {  
    "id": "USW00003192",  
    "country_code": "US",  
    "country": "United",  
    "location": {  
      "lat": 33.6228,  
      "lon": -111.9106  
    },  
    "elevation": 449,  
    "state_code": "AZ",  
    "state": "ARIZONA",  
    "name": "SCOTTSDALE MUNI AP"  
  },  
  "date": "2016-04-23T00:00:00",  
  "AWND": 2.4,  
  "PRCP": 0,  
  "TMAX": 29.4,  
  "TMIN": 18.9,  
}
```

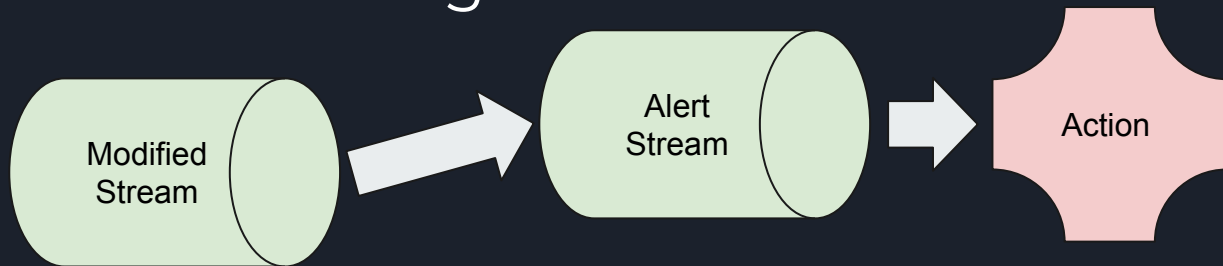
Say you had a farm near
USW00003192 and wanted
an email, SMS or
Notification if **TMIN** was
approaching freezing.



Weather Station Streaming Problem

Generalized

```
{
  "station": {
    "id": "USW00003192",
    "country_code": "US",
    "country": "United",
    "location": {
      "lat": 33.6228,
      "lon": -111.9106
    },
    "elevation": 449,
    "state_code": "AZ",
    "state": "ARIZONA",
    "name": "SCOTTSDALE MUNI AP"
  },
  "date": "2016-04-23T00:00:00",
  "AWND": 2.4,
  "PRCP": 0,
  "TMAX": 29.4,
  "TMIN": 18.9,
}
```

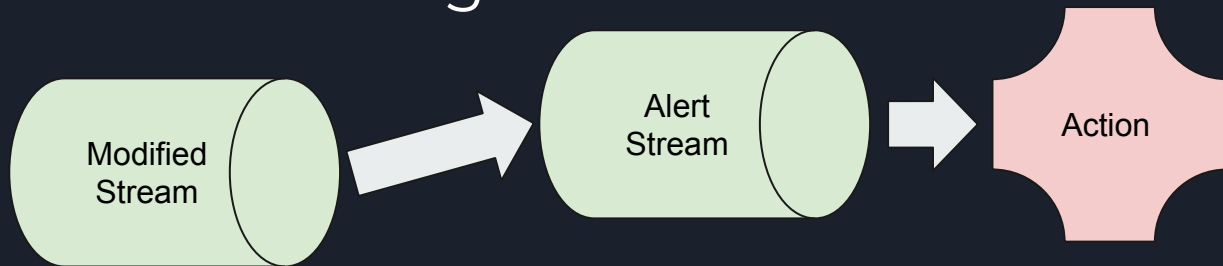


Filter for condition
and specific station
and alert user

Weather Station Streaming Problem

Generalized

```
{
  "station": {
    "id": "USW00003192",
    "country_code": "US",
    "country": "United",
    "location": {
      "lat": 33.6228,
      "lon": -111.9106
    },
    "elevation": 449,
    "state_code": "AZ",
    "state": "ARIZONA",
    "name": "SCOTTSDALE MUNI AP"
  },
  "date": "2016-04-23T00:00:00",
  "AWND": 2.4,
  "PRCP": 0,
  "TMAX": 29.4,
  "TMIN": 18.9,
}
```



Filter for condition
and specific station
and alert user

```
if (
  id === 'USW00003192'
  &&
  TMIN <= 1.0
)
```



Outline

- ETL and Stream Processing
- Teraslice Description and Features
- Teraslice Job
- Example Problem
- Teraslice Processor



Teraslice Asset/Processor Example

- Assets and Processors
 - A processor is the unit of code that manipulates the data
 - An asset is a collection of related processors
- Example Source
 - <https://github.com/godber/presentations/tree/main/phxjs-godber-teraslice-2023/weather-alert-asset>



Thank you!

Teraslice

Stream Processing with NodeJS

<https://github.com/terascope/teraslice>

<https://github.com/godber/presentations>

Austin Godber
@godber@az.social
May 17th, 2023